

# COMP7103C – Data Mining Assignment 2

<b>COMP7103C – Data Mining Assignment 2</b>	<b>1</b>
Group Information	1
Dataset Description	1
Dataset Preprocessing	4
Model Construction and Evaluation	5
Decision Tree	6
KNN	6
Logistic Regression	8
MLP	9
Naive Bayes	10
Random Forest	11
SVM	12
7-Output + MajorityVote	13
7-Output + MajorityVoteWithWeight	13
7-Output + MLP	13
7-Output + AdaBoost	14
GitHub	14

## Group Information

Name	UID
<b>ZHANG JIAYI</b>	3036032798
<b>Zheng Xufeng</b>	3036034019


Diagram 1 – Group Memebers

## Dataset Description

After receiving the dataset, we initially start dataset exploration. In this section, we visualized the datasets.



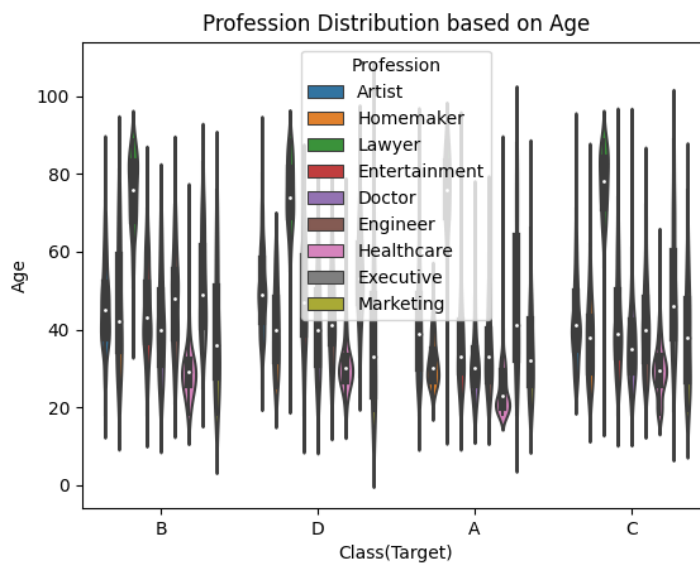
**Figure 1 – Class Distribution based on Gender**



**Figure 2 – Class Distribution based on Age**



**Figure 3 – Work Experience and Spending Score based on Class**



**Figure 4 – Profession Distribution based on Age**




Figure 5 – Correlation Matrix

## Dataset Preprocessing

We processed the dataset in the following manners.

1. Married: Fill all NaN with 0 if any.

```
df['Married'].fillna(0, inplace=True)
```

2. Gender: Fill all NaN with 2 if any.

```
df['Gender'].fillna(2, inplace=True)
```

3. Age: Fill all NaN with -1 if any.

```
df['Age'].fillna(-1, inplace=True)
```

4. Graduate: Fill all NaN with 2 if any.

```
df['Graduate'].fillna(2, inplace=True)
```

5. Profession: Replace 9 different professions with their corresponding integers, and fill all NaN with 0.

```
df.loc[df['Profession'].notna(), 'Profession'] =
df.loc[df['Profession'].notna(), 'Profession'].replace(
    {
        'Artist': 1,
        'Doctor': 2,
        'Engineer': 3,
        'Entertainment': 4,
        'Executive': 5,
        'Healthcare': 6,
        'Homemaker': 7,
        'Lawyer': 8,
        'Marketing': 9
    }
)
df['Profession'].fillna(0, inplace=True)
```

6. Years\_of\_Working : Fill all NaN with -1 if any.

```
df['Years_of_Working'].fillna(-1, inplace=True)
```

7. Spending\_Score: Replace 3 different spending scores with their corresponding integers, and fill all NaN with 0.

```
df.loc[df['Spending_Score'].notna(), 'Spending_Score'] =
df.loc[df['Spending_Score'].notna(), 'Spending_Score'].replace(
    {
        'Low': 1,
        'Average': 2,
        'High': 3
    })
df['Spending_Score'].fillna(0, inplace=True)
```

8. Family\_Members : Fill all NaN with 0 if any.

9. Category: Replace 7 different spending scores with their corresponding integers, and fill all NaN with 0.

```
df.loc[df['Category'].notna(), 'Category'] = df.loc[df['Category'].notna(),
'Category'].replace(
    {
        'Cat_1': 1,
        'Cat_2': 2,
        'Cat_3': 3,
        'Cat_4': 4,
        'Cat_5': 5,
        'Cat_6': 6,
        'Cat_7': 7
    })
df['Category'].fillna(0, inplace=True)
```

## Model Construction and Evaluation

In this section, we build 7 different models independently and also use 4 different ensemble learning methods to integrate the output of 7 different models' output. Hence there are totally 11 evaluated models.

Model	Accuracy on Validation Set	Output File Name
Decision Tree	<i>0.46960</i>	<i>Test_DT.csv</i>
KNN	<i>0.43966</i>	<i>Test_KNN.csv</i>
Logistic Regression	<i>0.44060</i>	<i>Test_LR.csv</i>
MLP	<i>0.45182</i>	<i>Test_MLP.csv</i>
Naive Bayes	<i>0.41628</i>	<i>Test_NB.csv</i>
Random Forest	<i>0.41534</i>	<i>Test_RF.csv</i>
SVM	<i>0.44808</i>	<i>Test_SVC.csv</i>

Diagram 2 – Weak Learning Evaluation Result

Model	Accuracy on Validation Set	Accuracy on Training Set	Output File Name
-------	----------------------------	--------------------------	------------------

7-Output + MajorityVote	0.4584	0.5119	<i>Test_Integration_MajorityVote.csv</i>
7-Output + MajorityVoteWithWeight	0.4153	0.9507	<i>Test_Integration_MajorityVoteWithWeight.csv</i>
7-Output + MLP	0.4153	0.9507	<i>Test_Integration_MLP.csv</i>
7-Output + AdaBoost	<b>0.4153</b>	<b>0.9507</b>	<i>Test_Integration_Ada.csv</i>

Diagram 3 – Ensemble Learning Evaluation Result

## Decision Tree

Build the model with the following parameters.

```
model = DecisionTreeClassifier(max_depth=4, criterion='entropy').fit(X_train, y_train)
```

The accuracy for the Validate set is 0.46960. Figure 6 shows in the confusion matrix.




Figure 6 – Confusion Matrix for Decision Tree on Validation Set

## KNN



First we find the tuned k among 100 candidates by using the following way.

```
iteration = 100
error_rate = []
acc = []
scores = {}
```

```

for i in range(1, iteration):
    model_knn = KNeighborsClassifier(n_neighbors=i)
    model_knn.fit(x_train, y_train)
    y_pred_knn = model_knn.predict(x_val)
    error_rate.append(np.mean(y_pred_knn != y_val))
    scores[i] = metrics.accuracy_score(y_val, y_pred_knn)
    acc.append(metrics.accuracy_score(y_val, y_pred_knn))

```




**Figure 7 – Different K Loss and Accuracy**

Yield that the best parameter for k is 18.

Build the model with the following parameters.

```
# TODO HERE
model_knn = KNeighborsClassifier(n_neighbors=18)
model_knn.fit(x_train, y_train)
```

The accuracy for the Validate set is 0.439663. Figure 8 shows in the confusion matrix.



**Figure 8 – Confusion Matrix for KNN on Validation Set**


## Logistic Regression

Build the model with the following parameters.

```
model = LogisticRegression(max_iter=600)
model.fit(x_train, y_train)
```

The accuracy for the Validate set is 0.440599. Figure 9 shows in the confusion matrix.






**Figure 9 – Confusion Matrix for Logistic Regression on Validation Set**

## MLP

Build the model with the following parameters.

```
model = MLPClassifier(solver='lbfgs', alpha=1e-5,
hidden_layer_sizes=(100,), random_state=1).fit(X_train, y_train)
```

The accuracy for the Validate set is 0.45182. Figure 10 shows in the confusion matrix.



**Figure 10 – Confusion Matrix for MLP on Validation Set**

## Naive Bayes

Build the model with the following parameters.

```
model = GaussianNB().fit(X_train, y_train)
```

The accuracy for the Validate set is 0.41628. Figure 11 shows in the confusion matrix.





Figure 11 – Confusion Matrix for MLP on Validation Set

## Random Forest

Build the model with the following parameters.

```
model_rf = RandomForestClassifier(n_estimators=100, criterion='gini',  
random_state=0)  
model_rf.fit(x_train, y_train)
```

The accuracy for the Validate set is 0.41534. Figure 12 shows in the confusion matrix.




**Figure 12 – Confusion Matrix for Random Forest on Validation Set**

## SVM

Build the model with the following parameters.

```
model = SVC(kernel='linear', C=1).fit(X_train, y_train)
```

The accuracy for the Validate set is 0.44808. Figure 13 shows in the confusion matrix.



**Figure 13 – Confusion Matrix for SVM on Validation Set**

### 7-Output + MajorityVote

We use the above 7 outputs together, and yield a majority results from those 7 outputs. The accuracy on the training set is 0.511954053693068, on the validating set is 0.4583723105706268.

```
df['majority_vote'] = df_train.apply(lambda x: x.mode()[0], axis=1)
df_s['majority_vote'] = df_val.apply(lambda x: x.mode()[0], axis=1)
```

### 7-Output + MajorityVoteWithWeight

Similar to above, the difference is we add 7-times weights on LR output. The accuracy on the training set is 0.9507145719246695, on the validating set is 0.4153414405986904.

### 7-Output + MLP

We took 7 output predictions as input for the MLP. The model is built as follows. We convert the labels into integers for the training and testing.

```
# Define the MLP model
model = Sequential()
model.add(Dense(64, input_dim=7, activation='relu'))
model.add(Dense(32, activation='relu'))
model.add(Dense(4, activation='softmax'))
# Compile the model
model.compile(loss='categorical_crossentropy', optimizer='adam',
metrics=['accuracy'])
```

The accuracy for the training set is 0.9507, on the validation set is 0.41534143686294556.

## 7-Output + AdaBoost

We took 7 output predictions as input for the AdaBoost. The model is built as follows. We convert the labels into integers for the training and testing.

```
model = AdaBoostClassifier(n_estimators=100, learning_rate=0.1,  
algorithm='SAMME')
```

The accuracy for the training set is 0.9507145719246695, on the validation set is 0.4153414405986904.

We eventually choose to use **AdaBoost** result as final results.

## GitHub

[https://github.com/AXNTROYUANXD/Simple\\_Classification](https://github.com/AXNTROYUANXD/Simple_Classification)